

15 – Awk / Gawk

CS 2043: Unix Tools and Scripting, Spring 2019 [1]

Matthew Milano

February 18, 2019

Cornell University

Table of Contents

1. AWK / GAWK
2. More about the filesystem

AWK / GAWK

awk Introduction

- **awk** is a programming language designed for processing text-based data.
 - Allows easy operation on fields rather than full lines.
 - Works in a *pattern-action* manner, like **sed**.
 - Supports numerical types (and operations).
 - Supports control-flow (e.g., **if** - **else** statements).
- Created at Bell Labs in the 1970s.
 - Alfred Aho, Peter Weinberger, and Brian Kernighan
 - An ancestor of **perl**, a *cousin* of **sed**.
 - Kernighan and Ritchie also invent C
- **Very** powerful.
 - It's *Turing Complete*!
 - ... a lot of things are.

gawk

- **gawk** is the GNU implementation of the **awk** programming language.
- On BSD/OSX, it is just called **awk**.
- On GNU, it is technically **gawk**, but should reliably be *symlinked* as **awk**.
- There are many different implementations of the AWK programming language.
 - If you use C or C++, this is similar to how there are different compilers. The compiler is an “implementation” of the language (big quotes on that...).
 - If you use Python, it’s like the difference between CPython, PyPy, Jython, etc.
 - Different implementations of the same programming language.

Basic Structure

- **awk** allows filters to handle text easily.
- The basic structure of an **awk** program is:

```
pattern1 { commands1 }  
pattern2 { commands2 }  
# ...
```

- Patterns can be regular expressions!
 - Proceeds line by line, checking each pattern one by one.
 - If the pattern is found, the { **commands** } are executed.
 - So for the above:
 - First line of input grabbed.
 - **pattern1** checked, if match { **commands1** } executed.
 - **pattern2** checked, if match { **commands2** } executed.
 - Next line of input grabbed.
 - Check **pattern1**, then **pattern2**, so on and so forth...

Why use **awk** over **sed**?

- Processing numerical values in **awk** is much more convenient.
- Variables and control flow in the actions.
- Convenient way of accessing fields within a given line.
- Flexible printing.
- Built-in arithmetic and string functions.
- Traditionally, **awk** has been used a lot in the scientific community e.g., biologists would use **awk** as a way of processing data and creating new table entries or something.
 - Basically, **awk** used to be the only real good *and* convenient option to process a large amount of data while still needing to perform mathematical computations or transformations.
 - These days there are many other options, but if you join a lab you may very well find some **awk** scripts creeping around and need to maintain them.

Simple Examples

- Print all lines containing Monster or monster.

```
awk '/[Mm]onster/ {print}' frankenstein.txt
```

- If no action specified, default is to print the whole line.

```
awk '/[Mm]onster/' frankenstein.txt
```

- The `$0` variable in `awk` refers to the whole line.

```
awk '/[Mm]onster/ {print $0}' frankenstein.txt
```

- First field (delimited by whitespace, or change *field separator*).

```
awk '/[Mm]onster/ {print $1}' frankenstein.txt
```

- `awk` understands extended regular expressions by default :)
 - We don't need to escape `+`, `?`, etc!

awk Shebang and BEGIN / END

- **awk** allows us blocks of code to be executed only once, at the beginning / end.
- With demo file **monstrosity.awk** and data file **frankenstein.txt** in current directory:

```
#!/usr/bin/awk -f
BEGIN { print "Starting search for monster..." }
/[Mm]onster/{ count++ } # Increment if [Mm]onster found
END { print "Found " count " monsters in the book." }
```

- Use the **-f** in the shebang to tell **awk** it expects a script.

```
$ ./monstrosity.awk # hangs... no input file
$ ./monstrosity.awk frankenstein.txt # yay!
# shebang '#!/usr/bin/awk -f' makes same as ...
$ awk -f monstrosity.awk frankenstein.txt
```

Using Variables in **awk**

- words are *variables by default*
 - opposite of **bash**, where words are strings by default
 - **word** is a variable (**\$word** works too)
- actions separated by semicolon
 - `{x = 0; y = 3; z = x + y; print z}`
- Not particularly whitespace sensitive!

Important Variables

- **NF**: the number of fields in the current line.
- **NR**: the number of lines read so far.
 - You cannot change **NF** or **NR**
- **FILENAME**: the name of the input file.
- **FS**: the *field separator*.
 - Example: change **FS** = `","` for processing a comma-separated-value sheet.
 - Can also specify `-F` flag (capital!) to set the **FS**.

Pattern Matching with **awk**

- **awk** can match any of the following pattern types:
 - **/regular expression/**
 - **relational expression**
 - **pattern1 && pattern2**
 - **pattern1 || pattern2**
 - **pattern1 ? pattern2: pattern3**
 - If **pattern1**, then match **pattern2**. Otherwise, match **pattern3**
 - **(pattern)**: parenthesis to group / change order of operations.
 - **! pattern** to invert **pattern**
 - **pattern1, pattern2**: match **pattern1**, work on every line until matches **pattern2**
 - So you cannot combine this...

Match action in (match) action

- there are *many* match-action programming languages.
 - sed
 - iptables
 - firewalls
 - datalog/prolog
- usually has *precedence*
 - take first match, like **case**.
- **awk** does **not** have precedence

Much Much More...

- Regular expression usage / comparison [available here](#).
- Many more comparison operations [detailed here](#).
- A wealth of useful / powerful built-in functions:
 - `toupper(x)`: make string upper case
 - `tolower(x)`: make string lower case
 - `exp(x)`: exponential of `x`
 - `rand()`: random number between `0` and `1`
 - `length(x)`: the length of `x`
 - `log(x)`: returns the log of `x`
 - `sin(x)`: returns the sine of `x`
 - `cos(x)`: returns the cosine of `x`
 - `int(x)`: convert
 - etc.
- Much more information [available here](#).

More about the filesystem

Inode the ultimate

- a data structure in a Unix-style file system that describes a file-system object such as a file or a directory.
- stores the attributes and disk block location(s) of the object's data.
- attributes may include metadata
 - (times of last change, access, modification)
 - owner and permission data.
- Directories are lists of names assigned to inodes.
- A directory contains an entry for itself, its parent, and each of its children.
- This was **all** cribbed straight from wikipedia. Go look!

Links in the filesystem

- When an **inode** for an object is in a directory, we say it's been **linked** into the filesystem tree
- the **ln** command makes and manages links.

link a filesystem object into the directory tree

```
ln [flags] <source> <target>
```

- works like **cp**; from **src** to **dst**
- creates a *peer* link; no notion of “original”
- only works on files

Symlinks in the filesystem

- A “soft” link or “symbolic” link isn’t a link at all
- works like a “shortcut” (really a **junction**) on Windows
- just a special file that contains a path in it
- looks light-blue under `ls`
 - you’ve seen this before!

symlink a filesystem object into the directory tree

```
ln -s [flags] <source> <target>
```

- technically the same command as `ln`, but used very differently with the `-s` flag!
- creates a *subordinate* link; refers to the path.
- doesn’t check to see if the source path was sensible first!
- works on files or directories.

References

- [1] Stephen McDowell, Bruno Abrahao, Hussam Abu-Libdeh, Nicolas Savva, David Slater, and others over the years. “Previous Cornell CS 2043 Course Slides”.