

15 – Networking and Package Management

CS 2043: Unix Tools and Scripting, Spring 2019 [1]

Matthew Milano

February 27, 2019

Cornell University

Table of Contents

1. welcome back to THE INTERNET
2. Package Management
3. System Specific Package Managers
4. Other Managers



Virtual Machines

CS2043 - Spring 2019
February 27

The image above is a link. Click it.

welcome back to THE INTERNET

Command we forgot from last time

ping a packet off a remote host

`ping [flags...] <host>`

- Simple echo back-and-forth
- tests connections
- uses **ICMP** protocol – same as **traceroute**
- runs forever by default

```
$ ping -c 4 google.com
PING google.com (172.217.9.238) 56(84) bytes of data.
64 bytes from lga34s11-in-f14.1e100.net (172.217.9.238): icmp_seq=1 ttl=55 time=8.24 ms
64 bytes from lga34s11-in-f14.1e100.net (172.217.9.238): icmp_seq=2 ttl=55 time=8.51 ms
64 bytes from lga34s11-in-f14.1e100.net (172.217.9.238): icmp_seq=3 ttl=55 time=8.56 ms
64 bytes from lga34s11-in-f14.1e100.net (172.217.9.238): icmp_seq=4 ttl=55 time=8.56 ms

--- google.com ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 8ms
rtt min/avg/max/mdev = 8.237/8.468/8.563/0.163 ms
```

Last time

- Computers communicate by sending **packets** through the network
- Packets are addressed to a local **MAC** and a potentially-remote **IP**
- **Switches** connect computers into a *local network* and forward packets by **MAC**
- **Routers** connect local networks into an *intranet* and forward packets by **IP**

Protocols from last time

- The **DHCP** protocol gives computers an IP address
- The **ARP** protocol associates an IP address with a MAC address
- The **DNS** protocol associates a domain name (google.com) with a MAC address

What is a protocol?

- an agreement on what sort of packets to exchange to achieve a particular goal
- Can be multi-step
- we distinguish between *transport layer* and *application layer*

More about protocols: transport layer

- *transport-layer* protocols correspond to different “kinds” of packets
 - examples: ARP, ICMP
- Operating system sees the different packets, handles them accordingly
- **operating system's job** to handle transport-layer packets

More about protocols: application layer

- *application-layer* protocols use the *same* kind of packet
 - examples: DHCP, DNS, HTTPS, SSH, most others you know
- Operating system passes them to applications
- How do applications find their packets?

Introducing: TCP and UDP

- *transport-layer* protocols for communicating with applications
- differentiate applications with “ports”
 - just a 16-bit integer
 - like apartment numbers
- applications listen at a specific port
 - registers with the OS
 - OS only forwards port-destined traffic
- contains “return addresses” for easy reply to client

- Most popular transport protocol
 - examples: HTTP, SSH
- *connection-oriented* protocol
 - “connect” to a port on a remote stream
 - receive a private channel on which to keep communicating
 - like a phone call ... or SSH session
- Hides common failures
 - ensures packets are reasonably ordered
 - retransmits packets if they get lost
 - cool algorithm to avoid congestion

UDP

- Second-most popular transport protocol
 - examples: DHCP, DNS, VoIP, Steam (as in video games), internet radio
 - **not** netflix
- **only gives you the port**
 - no connection: works like physical mail.
- All common failures exposed to application
 - packet order may vary
 - packets may not arrive
 - no indication whether transmitted packet got there
- Mostly used in either very-old, high-assurance or real-time applications
- more resilient to DOS attacks than TCP

Application protocols

- Still defines pattern of communication
- specific messages expected at specific times
- messages sent via (usually) TCP/UDP
- Example: HTTP, SSH, etc.

Exploring application protocols: netcat

netcat : so much more than **cat** over the **network**

```
nc [flags] [host]
```

```
nc -l -p <port>
```

```
nc <host> <port>
```

- Raw TCP protocol tool
- sends **stdin** over the network
- receives **stdout** from the network
- **nc -l** “listens”, behaves like a server
- **nc <host>** “connects”, behaves like a client

HTTP: a protocol to explore

- HTTP messages are raw text!
- Strings sent via TCP to port 80
- GET request: access a page

```
GET /people/mpmilano/ HTTP/1.1  
Host: cs.brown.edu
```

- Let's send this via **netcat**! (demo)
- Can explore more protocols this way; try it!

Some common ports

- HTTP: TCP/80
- SSH: TCP/22
- FTP: TCP/20 and TCP/21
- HTTPS: TCP/443
- SMTP (mail): TCP/25

Firewalls

- In a perfect world, we wouldn't need a firewall.
- Lives in the network, or in the kernel
- inspects traffic *before* it reaches its destination
- Two primary uses: filter legitimate services, block unwanted ones

Firewalls: the good uses

- Legit: *Filters* certain ports to prevent regions of the internet from accessing them
 - Cornell firewall drops all traffic destined to on-campus servers originating from off-campus IPs
 - **wash** firewall does the same
 - mail relay firewall would only allow known senders to connect
- prevents server from being overloaded by random external grievers
- prevents aggressive server scans from the darkweb
 - which, by the way, exists. ask me later.

Firewalls: the lazy uses.

- Block insecure / old apps
- cover up for weird/bad OS/system design
 - Example: print server on a mac at port 631
 - Example: just a lot of windows
- Block **all** uninvited remote connections
 - if your laptop isn't a server, shouldn't have exposed ports
 - if it does have exposed ports, some application is doing a bad.
- Fundamentally lazy: right answer is to secure the applications, not hide them.
- lots of legacy apps (that we're stuck with) can't be fixed, so also fundamentally necessary

Package Management

Package Management Overview

- If I had to give *only one reason* why Unix systems are superior to Windows: Package Management.
- Can install almost anything with ease of from your terminal.
- Update to the latest version with one command.
 - No more download the latest installer nonsense!
- Various tools can be installed by installing a *package*.
 - A package contains the files and other instructions to setup a piece of software.
 - Many packages depend on each other.
 - High-level package managers download packages, figure out the dependencies for you, and deal with groups of packages.
 - Low-level managers unpack individual packages, run scripts, and get the software installed correctly.
- In general, these are “pre-compiled binaries”: no compilation necessary. It’s already packaged nice and neat just for you!

Package Managers in the Wild

- GNU/Linux:
 - Low-level: two general families of *packages* exist: **deb**, and **rpm**.
 - High-level package managers you are likely to encounter:
 - Debian/Ubuntu: **apt-get**.
 - Some claim that **aptitude** is superior, but I will only cover **apt-get**. They are roughly interchangeable.
 - SUSE/OpenSUSE: **zypper**.
 - Fedora: **dnf** (Fedora 22+).
 - **zypper** and **dnf** use **SAT**-based dependency solvers, which many argue is fundamentally superior. The dependency resolution phase is usually not the slowest part though...installing the packages is. See [3] for more info.
 - RHEL/CentOS: **yum** (until they adopt **dnf**).
- Mac OSX:
 - Others exist, but the only one you should ever use is **brew**.
 - Don't user others (e.g. **port**), they are outdated / EOSL.

Using Package Managers

- Though the syntax for each package manager is different, the concepts are all the same.
 - This lecture will focus on **apt-get**, **dnf**, and **brew**.
 - The **dnf** commands are almost entirely interchangeable with **yum**, by design.
 - Note that **brew** is a “special snowflake”, more on this later.
- What does your package manager give you? The ability to
 - **install** new packages you do not have.
 - **remove** packages you have installed.
 - **update** installed packages.
 - update the lists to search for files / updates from.
 - view **dependencies** of a given package.
 - a whole lot more!!!

A Note on **update**

- The **update** command has importantly different meanings in different package managers.
- Some **do**, and some do **not** default to system (read linux kernel) updates.
 - Ubuntu: default is *no*.
 - Fedora: default is *yes*.
 - RHEL: default is *no*.
- It depends on your operating system, and package manager.
 - Know your operating system, and look up what the default behavior is.
- If your program needs a specific version of the linux kernel, you need to be very careful!

A Note on Names and their Meanings

- You may see packages of the form:
 - `<package>.i[3456]86` (e.g. `.i386` or `.i686`):
 - These are the **32-bit** packages.
 - `<package>.x86_64`: these are the **64-bit** packages.
 - `<package>.noarch`: these are independent of the architecture.
- Development tools can have as many as three packages:
 - The header files are usually called something like:
 - **deb**: usually `<package>-dev`
 - **rpm**: usually `<package>-devel`
 - The library you will need to link against:
 - If applicable, **lib**`<package>` or something similar.
 - The binaries (executables), often provided by just `<package>`.
 - Most relevant for **C** and **C++**, but also **Python** and others.
 - Use the **search** functionality of your package manager.

Example Development Tool Installation

- If I needed to compile and link against **Xrandr** (X.Org X11 libXrandr runtime library) on Fedora, I would have to install
 - **libXrandr**: the library.
 - **libXrandr-devel**: the header files.
 - Not including **.x86_64** is OK / encouraged, your package manager knows which one to install.
 - Though in certain special cases you may need to get the **32-bit** library as well.
 - In this case, if I were compiling a program that links against **libXrandr**, but I want to release a pre-compiled 32bit library, it must be installed in order for me to link against it.
- The **deb** versions should be similarly named, but just use the **search** functionality of find the right names.
- This concept has no meaning for **brew**, since it compiles everything.

System Specific Package Managers

Debian / Ubuntu Package Management (**apt-get**)

- Installing and uninstalling:
 - Install a package:
apt-get install <pkg1> <pkg2> ... <pkgN>
 - Remove a package:
apt-get remove <pkg1> <pkg2> ... <pkgN>
 - Only one **pkg** required, but can specify many.
 - “Group” packages are available, but still the same command.
- Updating components:
 - Update lists of packages available: **apt-get update**.
 - No arguments, it updates the whole list (even if you give args).
 - Updating currently installed packages: **apt-get upgrade**.
 - Specify a **package** name to only update / upgrade that package.
 - Update core (incl. kernel): **apt-get dist-upgrade**.
- Searching for packages:
 - Different command: **apt-cache search** <pkg>

RHEL / Fedora Package Managers (**yum** and **dnf**)

- Installing and uninstalling:
 - Install a package:
`dnf install <pkg1> <pkg2> ... <pkgN>`
 - Remove a package:
`dnf remove <pkg1> <pkg2> ... <pkgN>`
 - Only one **pkg** required, but can specify many.
 - “Group” packages are available, but different command:
 - `dnf groupinstall 'Package Group Name'`
- Updating components:
 - Update EVERYTHING: `dnf upgrade`.
 - **update** exists, but is essentially **upgrade**.
 - Specify a **package** name to only upgrade that package.
 - Updating repository lists: `dnf check-update`
- Searching for packages:
 - Same command: `dnf search <pkg>`
- **yum** and **dnf** (**Dandified Yum**) nearly interchangeable: [3].

dnf: Cautionary Tales

- **WARNING:** if you install package **Y**, which installs **X** as a dependency, and later **remove Y**
 - By default, **X** will be removed!
 - Refer to [2] for workarounds.
 - Generally, won't know you needed to **mark** until it is too late.
- Solution?
 - Basically, **pay attention to your package manager.**
 - It gets removed because nothing *explicitly* depends on it.
 - So one day you may realize "OH NO! I'm missing package **X**"...
 - ...so just **dnf install X**.
 - So while **mark** is available, personally I don't use it.
 - Sad face, I know. Just the way of the world.

OSX Package Management: Install **brew** on your own

- Sitting in class right now with a Mac?
- **DON'T DO THIS IN CLASS.** You will want to make sure you do not have to interrupt the process.
 - Make sure you have the “Command Line Tools” installed.
 - Instructions are on the [First Things First Config Page](#)
 - Visit <http://brew.sh/>
 - Copy-paste the given instructions in the terminal *as a regular user (not **root**.)*.
- **VERY IMPORTANT:** READ WHAT THE OUTPUT IS!!!! It will tell you to do things, and you *have* to do them. Specifically
You should run '**brew doctor**' BEFORE you install anything.

OSX Package Management (**brew**)

- Installing and uninstalling:
 - Install a *formula*:
`brew install <fmla1> <fmla2> ... <fmla2>`
 - Remove a formula:
`brew uninstall <fmla1> <fmla2> ... <fmlaN>`
 - Only one **fmla** required, but can specify many.
 - “Group” packages have no meaning in **brew**.
- Updating components:
 - Update **brew**, all *taps*, and installed formulae listings. This does not update the actual software you have installed with **brew**, just the definitions: `brew update`.
 - Update just installed formulae: `brew upgrade`.
 - Specify a **formula** name to only upgrade that formula.
- Searching for packages:
 - Same command: `brew search <formula>`

OSX: One of These Kids is Not Like the Others (Part I)

- Safe: confines itself (by default) in `/usr/local/Cellar`:
 - No **sudo**, plays nicely with OSX (e.g. Applications, **python3**).
 - Non-linking by default. If a conflict is detected, it will tell you.
 - **Really important to read what **brew** tells you!!!**
- **brew** is modular. Additional repositories ("*taps*") available:
 - Essentially what a **.rpm** or **.deb** would give you in linux.
 - These are 3rd party repos, not officially sanctioned by **brew**.
- Common taps people use:
 - **brew tap homebrew/science**
 - Various "scientific computing" tools, e.g. **opencv**.
 - **brew tap caskroom/cask**
 - Install **.app** applications! Safe: installs in the "Cellar", symlinks to `~/Applications`, but *now these update with brew all on their own* when you **brew update**!
 - E.g. **brew cask install vlc**

OSX: One of These Kids is Not Like the Others (Part II)

- **brew** installs *formulas*.
 - A **ruby** script that provides rules for where to download something from / how to compile it.
- Sometimes the packager creates a “**Bottle**”:
 - If a bottle for your version of OSX exists, you don’t have to compile locally.
 - The bottle just gets *downloaded* and then “*poured*”.
- Otherwise, **brew** downloads the source and compiles locally.
- Though more time consuming, can be quite convenient!
 - **brew options opencv**
 - **brew install --with-cuda --c++11 opencv**
 - It really really really is magical. No need to understand the **opencv** build flags, because the authors of the **brew** formula are kind and wonderful people.
 - **brew reinstall --with-missed-option formula**

OSX: One of These Kids is Not Like the Others (Part III)

- Reiteration: **pay attention to brew and what it says**. Seriously.
- Example: after installing **opencv**, it tells me:

==> Caveats

Python modules have been installed and Homebrews site-packages is not in your Python sys.path, so you will not be able to import the modules this formula installed. If you plan to develop with these modules, please run:

```
mkdir -p /Users/sven/.local/lib/python2.7/site-packages
echo 'import site; site.addsitedir(\
    "/usr/local/lib/python2.7/site-packages")' >> \
    /Users/sven/.local/lib/python2.7/site-packages/homebrew.pth
```

- **brew** gives copy-paste format, above is just so you can read.
- I want to use **opencv** in **Python**, so I do what **brew** tells me.

Less Common Package Management Operations

- Sometimes when dependencies are installed behind the scenes, and you no longer need them, you will want to get rid of them.
 - `apt-get autoremove`
 - `dnf autoremove`
 - `brew doctor`
- View the list of repositories being checked:
 - `apt-cache policy` (well, sort of...`apt` doesn't have it)
 - `dnf repolist [enabled|disabled|all]`
 - Some repositories for `dnf` are *disabled* by default (with good reason). Usually you want to just
`dnf --enablerepo=<name> install <thing>`
e.g. if you have `rawhide` (development branch for fedora).
 - `brew tap`

Other Managers

Like What?

- There are so many package managers out there for different things, too many to list them all!
- Ruby: **gem**
- Anaconda Python: **conda**
- Python: **pip**
- Python: **easy_install** (but really, just use **pip**)
- Python3: **pip3**
- LaTeX: **tlmgr** (uses the CTAN database)
 - Must install TeX from source to get **tlmgr**
- Perl: **cpan**
- Sublime Text: **Package Control**
- Many many others...

Like How?

- Some notes and warnings about Python package management.
- Notes:
 - If you want **X** in Python 2 **and** 3:
 - `pip install X` *and* `pip3 install X`
 - OSX Specifically: advise only using **brew** or Anaconda Python. The system Python can get really damaged if you modify it, you are better off leaving it alone.
 - So even if you want to use **python2** on Mac, I strongly encourage you to install it with **brew**.
- Warnings:
 - Don't mix **easy_install** and **pip**. Choose one, stick with it.
 - But the internet told me if I want **pip** on Mac, I should **easy_install pip**
 - NO! Because this **pip** will modify your **system** python, **USE BREW**.
 - Don't mix **pip** with **conda**. If you have Anaconda python, just stick to using **conda**.

References

- [1] Stephen McDowell, Bruno Abrahao, Hussam Abu-Libdeh, Nicolas Savva, David Slater, and others over the years. “Previous Cornell CS 2043 Course Slides”.
- [2] Reddit.com. *DNF Remove Package, keep dependencies??* 2016. URL: https://www.reddit.com/r/Fedora/comments/3pqrv9/dnf_remove_package_keep_dependencies/.
- [3] Jack Wallen. *What You Need to Know About Fedora’s Switch From Yum to DNF*. 2015. URL: <https://www.linux.com/learn/tutorials/838176-what-you-need-to-know-about-fedoras-switch-from-yum-to-dnf>.