

# 16 – Networking, OS, and Package Management

CS 2043: Unix Tools and Scripting, Spring 2019 [1]

---

Matthew Milano

March 1, 2019

Cornell University

# Table of Contents

1. Firewalls
2. Operating systems, and what they do.
3. Containers, and how they work
4. Package Management
5. System Specific Package Managers
6. Other Managers

# Firewalls

---

# Firewalls

- In a perfect world, we wouldn't need a firewall.
- Lives in the network, or in the kernel
- inspects traffic *before* it reaches its destination
- Two primary uses: filter legitimate services, block unwanted ones

# Firewalls: the good uses

- Legit: *Filters* certain ports to prevent regions of the internet from accessing them
  - Cornell firewall drops all traffic destined to on-campus servers originating from off-campus IPs
  - **wash** firewall does the same
  - mail relay firewall would only allow known senders to connect
- prevents server from being overloaded by random external grievers
- prevents aggressive server scans from the darkweb
  - which, by the way, exists. ask me later.

## Firewalls: the lazy uses.

- Block insecure / old apps
- cover up for weird/bad OS/system design
  - Example: print server on a mac at port 631
  - Example: just a lot of windows
- Block **all** uninvited remote connections
  - if your laptop isn't a server, shouldn't have exposed ports
  - if it does have exposed ports, some application is doing a bad.
- Fundamentally lazy: right answer is to secure the applications, not hide them.
- lots of legacy apps (that we're stuck with) can't be fixed, so also fundamentally necessary

Operating systems, and what they  
do.

---

# Processors

- The CPU; the chip at the center of your computer
- it actually runs your code
- wired via a *bus* to everything else in your computer
- Has multiple *cores* or *hyperthreads*
  - to allow code to execute simultaneously



# Processors have protection modes

- Pieces of code get associated with a protection mode
  - there's an instruction that literally says "when you run this code, drop these privileges"
- Protection modes let you drop lots of privileges
  - device access
  - physical memory access
  - ability to change protection modes
- Operating system always runs first and keeps all its privileges
- Operating system's job is to run processes for its users

# What is a process, really?

- A sequence of processor instructions
- runs from start to finish
- only thing running on CPU core
- what can a process do?
  - access its own memory
  - run arbitrary computation CPU commands
  - fire *interrupts*

# What is an interrupt?

- An “unexpected event”
- A request for something else to take over
- Like a signal (in C/unix), or Exception (in java/python/etc)
- Can register *interrupt handlers*, pieces of code that run interrupts
- The operating system registers itself as an *interrupt handler*
- A *syscall* is an interrupt handled by the OS
  - is how you read files, use network, etc.
  - OS registered the handler, so can have all privileges
  - most basic C functions / linux commands just fancy syscall wrappers!

# A potential process flow

- start a process
  - drop privileges
  - jump to process code
- do some computation
- read a file
  - fire an interrupt
  - interrupt handler (in OS) gets file
  - file placed in process memory
  - jump back to process code
- use file contents
- do more computation
- exit with result
  - fire an interrupt
  - interrupt handler (in OS) gets result
  - OS clears process memory

## Where VMs fit into this

- Using devices (from the OS) also interrupt-based!
- special instruction that sends message along system *bus*
- When host OS launches a VM
  - drops *some* privileges
  - registers *itself* (host OS) for device interrupts
  - launches guest OS
- when guest process wants to use a resource
  - interrupt back to guest OS
  - guest OS interrupts for device
    - Host OS gets interrupt
    - Host OS interrupts for device, or
    - Host OS takes over for a bit

# Containers, and how they work

---

# chrooting

change root directory

```
chroot <dir> <command>
```

- Must execute as root
- hides filesystem below <dir>
- **dir** looks like new /

- Why do this?
  - all **PATHs** relative to new root
  - system programs and libraries used from new root
  - can use programs that need incompatible libraries
  - can avoid upgrading system when using a program
- demo

- What's still the same in the chroot?
  - kernel
  - process space
  - RAM
  - devices
- Halfway to a container; can have a **chroot** of debian on ChromeOS
- No isolation between **chrooted** processes and “real” ones



# containers

- Special OS feature called LXC containers
- hides processes from each other
- can limit device access within a single container
  - how? checks PID after interrupt, denies request from container process
- 90% of a docker container is chroot + LXC
- Other 10%? Secretly a VM.
  - but only when needed
  - this is why “fancy” Windows 10 is required
- Docker build scripts and bundles are also nice

# Package Management

---

# Package Management Overview

- If I had to give *only one reason* why Unix systems are superior to Windows: Package Management.
- Can install almost anything with ease of from your terminal.
- Update to the latest version with one command.
  - No more download the latest installer nonsense!
- Various tools can be installed by installing a *package*.
  - A package contains the files and other instructions to setup a piece of software.
  - Many packages depend on each other.
  - High-level package managers download packages, figure out the dependencies for you, and deal with groups of packages.
  - Low-level managers unpack individual packages, run scripts, and get the software installed correctly.
- In general, these are “pre-compiled binaries”: no compilation necessary. It’s already packaged nice and neat just for you!

# Package Managers in the Wild

- GNU/Linux:
  - Low-level: two general families of *packages* exist: **deb**, and **rpm**.
  - High-level package managers you are likely to encounter:
    - Debian/Ubuntu: **apt-get**.
    - Some claim that **aptitude** is superior, but I will only cover **apt-get**. They are roughly interchangeable.
    - SUSE/OpenSUSE: **zypper**.
    - Fedora: **dnf** (Fedora 22+).
    - **zypper** and **dnf** use **SAT**-based dependency solvers, which many argue is fundamentally superior. The dependency resolution phase is usually not the slowest part though...installing the packages is. See [3] for more info.
    - RHEL/CentOS: **yum** (until they adopt **dnf**).
- Mac OSX:
  - Others exist, but the only one you should ever use is **brew**.
  - Don't user others (e.g. **port**), they are outdated / EOSL.

# Using Package Managers

- Though the syntax for each package manager is different, the concepts are all the same.
  - This lecture will focus on **apt-get**, **dnf**, and **brew**.
  - The **dnf** commands are almost entirely interchangeable with **yum**, by design.
  - Note that **brew** is a “special snowflake”, more on this later.
- What does your package manager give you? The ability to
  - **install** new packages you do not have.
  - **remove** packages you have installed.
  - **update** installed packages.
  - update the lists to search for files / updates from.
  - view **dependencies** of a given package.
  - a whole lot more!!!

## A Note on **update**

- The **update** command has importantly different meanings in different package managers.
- Some **do**, and some do **not** default to system (read linux kernel) updates.
  - Ubuntu: default is *no*.
  - Fedora: default is *yes*.
  - RHEL: default is *no*.
- It depends on your operating system, and package manager.
  - Know your operating system, and look up what the default behavior is.
- If your program needs a specific version of the linux kernel, you need to be very careful!

# A Note on Names and their Meanings

- You may see packages of the form:
  - `<package>.i[3456]86` (e.g. `.i386` or `.i686`):
    - These are the **32-bit** packages.
  - `<package>.x86_64`: these are the **64-bit** packages.
  - `<package>.noarch`: these are independent of the architecture.
- Development tools can have as many as three packages:
  - The header files are usually called something like:
    - **deb**: usually `<package>-dev`
    - **rpm**: usually `<package>-devel`
  - The library you will need to link against:
    - If applicable, **lib**`<package>` or something similar.
  - The binaries (executables), often provided by just `<package>`.
  - Most relevant for **C** and **C++**, but also **Python** and others.
  - Use the **search** functionality of your package manager.

## Example Development Tool Installation

- If I needed to compile and link against **Xrandr** (X.Org X11 libXrandr runtime library) on Fedora, I would have to install
  - **libXrandr**: the library.
  - **libXrandr-devel**: the header files.
  - Not including **.x86\_64** is OK / encouraged, your package manager knows which one to install.
  - Though in certain special cases you may need to get the **32-bit** library as well.
    - In this case, if I were compiling a program that links against **libXrandr**, but I want to release a pre-compiled 32bit library, it must be installed in order for me to link against it.
- The **deb** versions should be similarly named, but just use the **search** functionality of find the right names.
- This concept has no meaning for **brew**, since it compiles everything.



# System Specific Package Managers

---

# Debian / Ubuntu Package Management (**apt-get**)

- Installing and uninstalling:
  - Install a package:  
**apt-get install** <pkg1> <pkg2> ... <pkgN>
  - Remove a package:  
**apt-get remove** <pkg1> <pkg2> ... <pkgN>
  - Only one **pkg** required, but can specify many.
  - “Group” packages are available, but still the same command.
- Updating components:
  - Update lists of packages available: **apt-get update**.
    - No arguments, it updates the whole list (even if you give args).
  - Updating currently installed packages: **apt-get upgrade**.
    - Specify a **package** name to only update / upgrade that package.
  - Update core (incl. kernel): **apt-get dist-upgrade**.
- Searching for packages:
  - Different command: **apt-cache search** <pkg>

## RHEL / Fedora Package Managers (**yum** and **dnf**)

- Installing and uninstalling:
  - Install a package:  
`dnf install <pkg1> <pkg2> ... <pkgN>`
  - Remove a package:  
`dnf remove <pkg1> <pkg2> ... <pkgN>`
  - Only one **pkg** required, but can specify many.
  - “Group” packages are available, but different command:
    - `dnf groupinstall 'Package Group Name'`
- Updating components:
  - Update EVERYTHING: `dnf upgrade`.
  - **update** exists, but is essentially **upgrade**.
    - Specify a **package** name to only upgrade that package.
  - Updating repository lists: `dnf check-update`
- Searching for packages:
  - Same command: `dnf search <pkg>`
- **yum** and **dnf** (**Dandified Yum**) nearly interchangeable: [3].

## dnf: Cautionary Tales

- **WARNING:** if you install package **Y**, which installs **X** as a dependency, and later **remove Y**
  - By default, **X** will be removed!
  - Refer to [2] for workarounds.
  - Generally, won't know you needed to **mark** until it is too late.
- Solution?
  - Basically, **pay attention to your package manager.**
  - It gets removed because nothing *explicitly* depends on it.
  - So one day you may realize "OH NO! I'm missing package **X**"...
  - ...so just **dnf install X**.
    - So while **mark** is available, personally I don't use it.
  - Sad face, I know. Just the way of the world.

# OSX Package Management: Install **brew** on your own

- Sitting in class right now with a Mac?
- **DON'T DO THIS IN CLASS.** You will want to make sure you do not have to interrupt the process.
  - Make sure you have the “Command Line Tools” installed.
    - Instructions are on the [First Things First Config Page](#)
  - Visit <http://brew.sh/>
  - Copy-paste the given instructions in the terminal *as a regular user (not **root**.)*.
- **VERY IMPORTANT:** READ WHAT THE OUTPUT IS!!!! It will tell you to do things, and you *have* to do them. Specifically  
You should run '**brew doctor**' BEFORE you install anything.

# OSX Package Management (**brew**)

- Installing and uninstalling:
  - Install a *formula*:  
`brew install <fmla1> <fmla2> ... <fmla2>`
  - Remove a formula:  
`brew uninstall <fmla1> <fmla2> ... <fmlaN>`
  - Only one **fmla** required, but can specify many.
  - “Group” packages have no meaning in **brew**.
- Updating components:
  - Update **brew**, all *taps*, and installed formulae listings. This does not update the actual software you have installed with **brew**, just the definitions: `brew update`.
  - Update just installed formulae: `brew upgrade`.
    - Specify a **formula** name to only upgrade that formula.
- Searching for packages:
  - Same command: `brew search <formula>`

## OSX: One of These Kids is Not Like the Others (Part I)

- Safe: confines itself (by default) in `/usr/local/Cellar`:
  - No **sudo**, plays nicely with OSX (e.g. Applications, **python3**).
  - Non-linking by default. If a conflict is detected, it will tell you.
  - **Really important to read what **brew** tells you!!!**
- **brew** is modular. Additional repositories (“*taps*”) available:
  - Essentially what a **.rpm** or **.deb** would give you in linux.
  - These are 3rd party repos, not officially sanctioned by **brew**.
- Common taps people use:
  - **brew tap homebrew/science**
    - Various “scientific computing” tools, e.g. **opencv**.
  - **brew tap caskroom/cask**
    - Install **.app** applications! Safe: installs in the “Cellar”, symlinks to `~/Applications`, but *now these update with brew all on their own* when you **brew update**!
    - E.g. **brew cask install vlc**

## OSX: One of These Kids is Not Like the Others (Part II)

- **brew** installs *formulas*.
  - A **ruby** script that provides rules for where to download something from / how to compile it.
- Sometimes the packager creates a “**Bottle**”:
  - If a bottle for your version of OSX exists, you don't have to compile locally.
  - The bottle just gets *downloaded* and then “*poured*”.
- Otherwise, **brew** downloads the source and compiles locally.
- Though more time consuming, can be quite convenient!
  - **brew options opencv**
  - **brew install --with-cuda --c++11 opencv**
  - It really really really is magical. No need to understand the **opencv** build flags, because the authors of the **brew** formula are kind and wonderful people.
  - **brew reinstall --with-missed-option formula**



## OSX: One of These Kids is Not Like the Others (Part III)

- Reiteration: **pay attention to brew and what it says**. Seriously.
- Example: after installing **opencv**, it tells me:

**==> Caveats**

Python modules have been installed and Homebrews site-packages is not in your Python sys.path, so you will not be able to import the modules this formula installed. If you plan to develop with these modules, please run:

```
mkdir -p /Users/sven/.local/lib/python2.7/site-packages
echo 'import site; site.addsitedir(
    "/usr/local/lib/python2.7/site-packages")' >> \
    /Users/sven/.local/lib/python2.7/site-packages/homebrew.pth
```

- **brew** gives copy-paste format, above is just so you can read.
- I want to use **opencv** in **Python**, so I do what **brew** tells me.

## Less Common Package Management Operations

- Sometimes when dependencies are installed behind the scenes, and you no longer need them, you will want to get rid of them.
  - `apt-get autoremove`
  - `dnf autoremove`
  - `brew doctor`
- View the list of repositories being checked:
  - `apt-cache policy` (well, sort of...`apt` doesn't have it)
  - `dnf repolist [enabled|disabled|all]`
    - Some repositories for `dnf` are *disabled* by default (with good reason). Usually you want to just  
`dnf --enablerepo=<name> install <thing>`  
e.g. if you have `rawhide` (development branch for fedora).
  - `brew tap`

## Other Managers

---

## Like What?

- There are so many package managers out there for different things, too many to list them all!
- Ruby: **gem**
- Anaconda Python: **conda**
- Python: **pip**
- Python: **easy\_install** (but really, just use **pip**)
- Python3: **pip3**
- LaTeX: **tlmgr** (uses the CTAN database)
  - Must install TeX from source to get **tlmgr**
- Perl: **cpan**
- Sublime Text: **Package Control**
- Many many others...

## Like How?

- Some notes and warnings about Python package management.
- Notes:
  - If you want **X** in Python 2 **and** 3:
    - `pip install X` *and* `pip3 install X`
  - OSX Specifically: advise only using **brew** or Anaconda Python. The system Python can get really damaged if you modify it, you are better off leaving it alone.
  - So even if you want to use **python2** on Mac, I strongly encourage you to install it with **brew**.
- Warnings:
  - Don't mix **easy\_install** and **pip**. Choose one, stick with it.
    - But the internet told me if I want **pip** on Mac, I should **easy\_install pip**
    - NO! Because this **pip** will modify your **system** python, **USE BREW**.
  - Don't mix **pip** with **conda**. If you have Anaconda python, just stick to using **conda**.

# References

- [1] Stephen McDowell, Bruno Abrahao, Hussam Abu-Libdeh, Nicolas Savva, David Slater, and others over the years. “Previous Cornell CS 2043 Course Slides”.
- [2] Reddit.com. *DNF Remove Package, keep dependencies??* 2016. URL: [https://www.reddit.com/r/Fedora/comments/3pqr99/dnf\\_remove\\_package\\_keep\\_dependencies/](https://www.reddit.com/r/Fedora/comments/3pqr99/dnf_remove_package_keep_dependencies/).
- [3] Jack Wallen. *What You Need to Know About Fedora’s Switch From Yum to DNF*. 2015. URL: <https://www.linux.com/learn/tutorials/838176-what-you-need-to-know-about-fedoras-switch-from-yum-to-dnf>.